

Definitive Implementation Guide: A High-Performance KV Store

This document provides a complete roadmap for developing the high-performance transactional key-value store. It includes a full macOS development environment setup, a detailed project and directory structure, and a granular, phased implementation plan with specific tasks, deliverables, and verification steps.

1. Development Environment Setup (macOS)

This section details the tools and steps required to configure a C++ development environment on macOS.

1.1. Core Tools Installation

1. **Xcode Command Line Tools:** This package provides the Clang C++ compiler, Git, and other essential command-line utilities.
 - **Action:** Open the Terminal and run:

```
Bash  
xcode-select --install
```
 - **Verification:** After installation, verify that the Clang compiler is available:

```
Bash  
clang --version
```
2. **Homebrew:** A package manager for macOS used to install software not provided by Apple.
 - **Action:** Install Homebrew by running the command from its official website:

```
Bash  
/bin/bash -c "$(curl -fsSL  
https://raw.githubusercontent.com/Homebrew/install/HEAD/install.sh)"
```

3. **CMake:** The build system generator for managing the project's compilation process.

- **Action:** Use Homebrew to install CMake ¹:

```
Bash  
brew install cmake
```

- **Verification:** Check the installed version:

```
Bash  
cmake --version
```

1.2. Editor and Extensions

- **Visual Studio Code:** A lightweight but powerful code editor.
 - **Action:** Download and install from the official website.
 - **C/C++ Extension Pack:** For debugging, code completion, and syntax highlighting.
 - **Action:** Inside VS Code, open the Extensions view (Cmd+Shift+X), search for C/C++ Extension Pack, and install it.¹

2. Project Directory and File Structure

A well-organized directory structure is crucial for maintainability and scalability. The following structure separates concerns such as source code, public headers, tests, and dependencies.²

```
kv_store/  
├──.gitignore  
├──CMakeLists.txt      # Root CMake file: project definition, enables testing, adds  
subdirectories  
├── build/             # Build artifacts (ignored by Git)  
├── docs/              # Project documentation (design docs, etc.)  
├── include/  
└── kv_store/          # Public headers, namespaced by project name
```

```

|   |   |— common/
|   |   |   |— bloom_filter.h
|   |   |   |— status.h
|   |   |— io/
|   |   |   |— file_manager.h
|   |   |   |— sstable.h
|   |   |— memtable/
|   |   |   |— memtable.h
|   |   |— transaction/
|   |   |   |— transaction.h
|   |   |— wal/
|   |   |   |— wal.h
|   |— src/
|   |   |— CMakeLists.txt # Builds the main static library `libkv_store.a`
|   |   |— common/
|   |   |   |— bloom_filter.cpp
|   |   |   |— status.cpp
|   |   |— io/
|   |   |   |— file_manager.cpp
|   |   |   |— sstable_builder.cpp
|   |   |   |— sstable_reader.cpp
|   |   |   |— sstable_format.h # Private header for SSTable block structures
|   |   |— index/
|   |   |   |— art.cpp
|   |   |   |— art.h # Private header for ART implementation
|   |   |— memtable/
|   |   |   |— memtable.cpp
|   |   |   |— skip_list.h # Private header for lock-free Skip List
|   |   |— transaction/
|   |   |   |— transaction.cpp
|   |   |   |— mvcc.cpp
|   |   |— wal/
|   |   |   |— wal_reader.cpp
|   |   |   |— wal_writer.cpp
|   |   |   |— wal_format.h # Private header for WAL record structure
|   |— tests/
|   |   |— CMakeLists.txt # Defines test executables, links against GTest and our library
|   |   |— common/
|   |   |   |— test_bloom_filter.cpp
|   |   |   |— test_status.cpp
|   |   |— io/
|   |   |   |— test_sstable.cpp
|   |   |— memtable/

```

```
|   └─ test_memtable.cpp
|   └─ transaction/
|       └─ test_mvcc.cpp
|       └─ wal/
|           └─ test_wal.cpp
```

3. Granular Implementation Plan

This plan breaks down the development process into four distinct phases. Each task is designed to be a small, verifiable step.

Phase 1: Core Data Structures & I/O Primitives

Goal: Build the fundamental, self-contained building blocks. All components in this phase can be developed and unit-tested in isolation.

Module 1.1: Low-Level I/O Abstraction

- **Task 1.1.1: Define File System Interfaces**
 - **Description:** Create abstract base classes `IWritableFile`, `IReadableFile`, and `IFileManager` to define a clean interface for file operations.
 - **Deliverable:** `include/kv_store/io/file_manager.h`.
- **Task 1.1.2: Implement POSIX File System Wrapper**
 - **Description:** Implement concrete classes that use standard C++ file streams or POSIX system calls (`open`, `read`, `write`, `fsync`).⁵
 - **Deliverable:** `src/io/file_manager.cpp`.
- **Task 1.1.3: Create Unit Tests with Mocking**
 - **Description:** Write unit tests for the I/O layer. Use a mocking framework (like Google Mock, part of Google Test) to simulate disk behavior.
 - **Deliverable:** `tests/io/test_file_manager.cpp`.
 - **Verification:** Confirm all interface methods behave as expected under normal and error conditions.

Module 1.2: On-Disk Data Formats (Serialization)

- **Task 1.2.1: Implement WAL Record Serialization**
 - **Description:** Define the WALRecord struct in a private header. Implement serialize and deserialize methods. The format must be ``.⁵
 - **Deliverable:** src/wal/wal_format.h.
 - **Verification:** Unit test in tests/wal/test_wal.cpp that a serialized record can be deserialized perfectly and that corruption is detected.
- **Task 1.2.2: Implement SSTable Block Formats**
 - **Description:** In a private header, define structs for BlockHandle (offset, size), and SSTableFooter (magic_number, index_block_handle, filter_block_handle).
 - **Deliverable:** src/io/sstable_format.h.
- **Task 1.2.3: Implement Data Block Builder & Reader**
 - **Description:** Create a DataBlockBuilder to add key-value pairs and generate a serialized data block with prefix compression. Create a DataBlockReader to parse it.
 - **Deliverable:** Add classes to src/io/sstable_builder.cpp and src/io/sstable_reader.cpp.
 - **Verification:** Unit test in tests/io/test_sstable.cpp by building a block and verifying the reader can retrieve all original key-value pairs.
- **Task 1.2.4: Implement Index Block Builder & Reader**
 - **Description:** Create an IndexBlockBuilder that takes divider keys and BlockHandles, and an IndexBlockReader that can perform a binary search.
 - **Deliverable:** Add classes to src/io/sstable_builder.cpp and src/io/sstable_reader.cpp.
 - **Verification:** Unit test in tests/io/test_sstable.cpp by building an index block and verifying lookups.

Module 1.3: Standalone Data Structures (Single-Threaded)

- **Task 1.3.1: Implement Bloom Filter**
 - **Description:** Create a BloomFilter class with add(key) and might_contain(key) methods. It must be serializable to a byte buffer.
 - **Deliverable:** include/kv_store/common/bloom_filter.h and src/common/bloom_filter.cpp.
 - **Verification:** Unit test in tests/common/test_bloom_filter.cpp. Verify no false negatives and that the false positive rate is within configured bounds.
- **Task 1.3.2: Implement Skip List**
 - **Description:** Implement a standard, single-threaded Skip List for storing sorted key-value pairs.

- **Deliverable:** A SkipList class in the private header `src/memtable/skip_list.h`.
 - **Verification:** Unit test in `tests/memtable/test_memtable.cpp`. Verify all operations and ordered iteration.
 - **Task 1.3.3: Implement Adaptive Radix Tree (ART)**
 - **Description:** Implement a single-threaded ART with adaptive node types (Node4, Node16, Node48, Node256) and path compression.
 - **Deliverable:** An ART class in the private header `src/index/art.h` and implementation in `src/index/art.cpp`.
 - **Verification:** Write extensive unit tests in `tests/index/test_art.cpp` covering all operations and node type transitions.
-

Phase 2: Building the Write & Read Path

Goal: Assemble the Phase 1 components into functional, single-threaded data paths.

Module 2.1: WAL Manager

- **Task 2.1.1: Implement WALWriter**
 - **Dependencies:** I/O Abstraction (1.1), WAL Record Format (1.2.1).
 - **Description:** Create a class to manage writing to a log file. Implement group commit logic to buffer records and sync them in a single batch.⁷
 - **Deliverable:** WALWriter class in `src/wal/wal_writer.cpp`.
- **Task 2.1.2: Implement WALReader and Recovery**
 - **Dependencies:** I/O Abstraction (1.1), WAL Record Format (1.2.1).
 - **Description:** Create a WALReader to read records sequentially, validating checksums. Implement a `recover()` function that replays log entries.
 - **Deliverable:** WALReader class in `src/wal/wal_reader.cpp`.
- **Task 2.1.3: Test WAL Durability**
 - **Dependencies:** WALWriter, WALReader.
 - **Description:** In `tests/wal/test_wal.cpp`, write a test that writes records, simulates a crash (by not closing the file), and verifies `recover()` reconstructs the state.

Module 2.2: SSTable Manager

- **Task 2.2.1: Implement SSTableBuilder**
 - **Dependencies:** I/O Abstraction (1.1), All SSTable Formats (1.2), Bloom Filter (1.3.1).
 - **Description:** Create a class that orchestrates writing a complete SSTable file.
 - **Deliverable:** SSTableBuilder class in src/io/sstable_builder.cpp.
- **Task 2.2.2: Implement SSTableReader**
 - **Dependencies:** I/O Abstraction (1.1), All SSTable Formats (1.2), Bloom Filter (1.3.1).
 - **Description:** Create a class to open an SSTable, load its index and filter, and perform key lookups.
 - **Deliverable:** SSTableReader class in src/io/sstable_reader.cpp.
- **Task 2.2.3: Test SSTable Read/Write Path**
 - **Dependencies:** SSTableBuilder, SSTableReader.
 - **Description:** In tests/io/test_sstable.cpp, write an integration test that builds an SSTable, then uses the reader to perform lookups.
 - **Verification:** Confirm correct values are returned and the Bloom filter prevents disk reads for non-existent keys.

Module 2.3: Memtable and Flush Logic

- **Task 2.3.1: Create Memtable Class**
 - **Dependencies:** Skip List (1.3.2).
 - **Description:** Create a Memtable class that wraps the SkipList and tracks its memory usage.
 - **Deliverable:** include/kv_store/memtable/memtable.h and src/memtable/memtable.cpp.
- **Task 2.3.2: Implement Flush Process**
 - **Dependencies:** Memtable (2.3.1), SSTableBuilder (2.2.1).
 - **Description:** Implement a method in the Memtable class that takes its sorted iterator and feeds the key-value pairs into an SSTableBuilder.
 - **Deliverable:** flush_to_sstable method in src/memtable/memtable.cpp.
- **Task 2.3.3: Test Memtable Flush**
 - **Dependencies:** Memtable (2.3.1), SSTableReader (2.2.2).
 - **Description:** In tests/memtable/test_memtable.cpp, populate a Memtable, flush it, and use an SSTableReader to verify the contents.

Phase 3: Concurrency and Transactions

Goal: Introduce the MVCC framework and make the data paths thread-safe and transactional.

Module 3.1: Lock-Free Data Structures

- **Task 3.1.1: Implement ConcurrentSkipList**
 - **Dependencies:** Single-threaded Skip List (1.3.2).
 - **Description:** Re-implement the SkipList in `src/memtable/skip_list.h` using `std::atomic` pointers and compare-and-swap (CAS) loops for thread-safe operations.⁸
 - **Deliverable:** A thread-safe `ConcurrentSkipList` class.
 - **Verification:** Write multi-threaded stress tests in `tests/memtable/test_memtable.cpp`. Use tools like ThreadSanitizer to detect data races.
- **Task 3.1.2: Implement Persistent ART with Path Copying**
 - **Dependencies:** Single-threaded ART (1.3.3).
 - **Description:** Modify the ART's update methods to be non-destructive, using path copying to create a new tree version and return a new root pointer.
 - **Deliverable:** A persistent (functional) ART class in `src/index/art.h` and `src/index/art.cpp`.
 - **Verification:** Unit test that an update creates a new root, and that traversals from old and new roots see the correct data.

Module 3.2: MVCC Core

- **Task 3.2.1: Implement TransactionManager**
 - **Description:** Create a manager that issues unique TxnIDs using `std::atomic<uint64_t>` and tracks active transactions thread-safely.
 - **Deliverable:** `TransactionManager` class in `src/transaction/mvcc.cpp`.
- **Task 3.2.2: Implement Versioned Data Records**
 - **Description:** Modify all data-holding structures to store versioned records: (key, creation_TxnID, deletion_TxnID, value). A delete operation inserts a record with a "tombstone" marker.
 - **Deliverable:** Updated data record definitions in private headers.
- **Task 3.2.3: Implement Transaction Snapshots and Visibility**
 - **Description:** Create a `Transaction` class that gets a snapshot from the `TransactionManager`. Implement the core `isVisible(record, snapshot)` function according to the rules of Snapshot Isolation.
 - **Deliverable:** `include/kv_store/transaction/transaction.h`, `src/transaction/transaction.cpp`, and the `isVisible` function in

src/transaction/mvcc.cpp.

- **Verification:** Write extensive unit tests for the isVisible function in tests/transaction/test_mvcc.cpp.

Module 3.3: Persistent & Versioned ART Integration

- **Task 3.3.1: Implement ART Persistence and Lazy Loading**

- **Dependencies:** Persistent ART (3.1.2).
 - **Description:** Implement the on-disk format using a post-order traversal and the pointer swizzling mechanism using the most significant bit of a 64-bit pointer as a flag.
 - **Deliverable:** Serialization and lazy-loading logic for the ART in src/index/art.cpp.
 - **Verification:** Create a large ART, serialize it, restart, and perform lookups. Verify that only necessary nodes are read from disk.
-

Phase 4: Background Processes & Final Integration

Goal: Implement maintenance tasks and assemble the final, user-facing database engine.

Module 4.1: Compaction Engine

- **Task 4.1.1: Implement MANIFEST Logic**

- **Description:** Implement classes to represent the LSM-tree state (Version, VersionEdit). Implement logic to write a new MANIFEST file and atomically update the database state using a rename on a CURRENT file.
- **Deliverable:** Version management and MANIFEST I/O classes in src/io/version.cpp (and private header).

- **Task 4.1.2: Implement K-Way Merge Iterator**

- **Description:** Create a merging iterator that takes a list of SSTable iterators and yields key-value pairs in sorted order, correctly handling multiple versions of the same key.
- **Deliverable:** A generic merging iterator in src/io/merge_iterator.h.

- **Task 4.1.3: Implement Leveled Compaction Strategy**

- **Dependencies:** MANIFEST (4.1.1), Merge Iterator (4.1.2), SSTable Manager (2.2).
- **Description:** Implement the logic for Leveled Compaction: decide when and what to compact.

- **Deliverable:** Compaction logic in src/compaction/compactor.cpp.
- **Task 4.1.4: Implement Background Thread Pool**
 - **Description:** Create a dedicated thread pool for running flush and compaction jobs.
 - **Deliverable:** A configurable background thread pool in src/common/thread_pool.cpp.

Module 4.2: Garbage Collection

- **Task 4.2.1: Integrate Data GC with Compaction**
 - **Dependencies:** Compaction Engine (4.1), TransactionManager (3.2.1).
 - **Description:** The TransactionManager must expose a getLowWatermark() method. The compactor must use this to discard any version whose deletion_TxnID is committed and older than the low-water mark.
 - **Deliverable:** GC logic integrated into src/compaction/compactor.cpp.
- **Task 4.2.2: Implement ART Node GC**
 - **Dependencies:** Persistent ART (3.1.2).
 - **Description:** Implement a mark-and-sweep garbage collector for ART nodes. Roots are the latest committed ART root and all roots held by active transaction snapshots.
 - **Deliverable:** A garbage collector for persistent ART nodes in src/index/art_gc.cpp.
- **Task 4.2.3: Test Garbage Collection**
 - **Description:** Write a test that runs a workload with updates/deletes, holds a long-running transaction to pin the low-water mark, runs compaction, and verifies old versions are kept. Then, commit the transaction, advance the mark, run compaction again, and verify the old versions are now purged.

Module 4.3: Top-Level API Integration & Build System

- **Task 4.3.1: Implement Database Class**
 - **Dependencies:** All other components.
 - **Description:** Create the final top-level class that encapsulates all components and exposes the public API: Open, Close, BeginTransaction.
 - **Deliverable:** The main Database class in src/db.cpp.
- **Task 4.3.2: Configure Root CMakeLists.txt**
 - **Description:** Set up the main project file.
 - **Deliverable:** CMakeLists.txt with:

```
CMake
cmake_minimum_required(VERSION 3.14)
project(kv_store CXX)
```

```
set(CMAKE_CXX_STANDARD 17)
set(CMAKE_CXX_STANDARD_REQUIRED ON)
```

```
# Enable testing
enable_testing()
```

```
# Add subdirectories
add_subdirectory(src)
add_subdirectory(tests)
```

- **Task 4.3.3: Configure src/CMakeLists.txt**

- **Description:** Define the main library target.
- **Deliverable:** src/CMakeLists.txt with:

```
CMake
# Gather all source files
file(GLOB_RECURSE LIB_SOURCES "*.cpp")

# Create the static library
add_library(kv_store STATIC ${LIB_SOURCES})

# Expose the public include directory
target_include_directories(kv_store PUBLIC ../include)
```

- **Task 4.3.4: Configure tests/CMakeLists.txt**

- **Description:** Set up Google Test and the test executable.¹²
- **Deliverable:** tests/CMakeLists.txt with:

```
CMake
# Download and configure Google Test
include(FetchContent)
FetchContent_Declare(googletest URL
https://github.com/google/googletest/archive/03597a01ee50ed33e9dfd640b249b4be3799d395.zip)
FetchContent_MakeAvailable(googletest)

# Gather all test source files
file(GLOB_RECURSE TEST_SOURCES "*.cpp")

# Create the test executable
add_executable(run_tests ${TEST_SOURCES})

# Link against our library and GTest
target_link_libraries(run_tests PRIVATE kv_store GTest::gtest_main)

# Discover tests with CTest
```

```
include(GoogleTest)
gtest_discover_tests(run_tests)
```

- **Task 4.3.5: Write End-to-End Integration and Stress Tests**

- **Description:** Create comprehensive tests that simulate real-world application workloads, including concurrent transactions, high write volumes, and crash-recovery scenarios.
- **Verification:** The database must remain consistent and correct under high concurrent load and after simulated crashes.

Works cited

1. Setup Mac OS for C++, CMake development using Visual Studio Code - Medium, accessed September 4, 2025, <https://medium.com/@divyendu.narayan/visual-studio-code-setup-in-mac-os-to-build-and-debug-c-cmake-projects-45a78b29e49>
2. How to Structure C Projects: These Best Practices Worked for Me | Blog - Luca Cavallin, accessed September 4, 2025, <https://www.lucavall.in/blog/how-to-structure-c-projects-my-experience-best-practices>
3. Best Practices for Structuring C++ Projects (Industry Standards) | by Govind Yadav - Medium, accessed September 4, 2025, <https://medium.com/@gtech.govind2000/best-practices-for-structuring-c-projects-industry-standards-71b82f6b145c>
4. Structuring a C++ Project Effectively - slaptijack, accessed September 4, 2025, <https://slaptijack.com/programming/project-structure-for-cpp.html>
5. Filesystem Home - Boost, accessed September 4, 2025, <https://www.boost.org/libs/filesystem/>
6. Filesystem library (since C++17) - cppreference.com - C++ Reference, accessed September 4, 2025, <https://en.cppreference.com/w/cpp/filesystem.html>
7. Modern C++ Project Structuring | Gregory Kelleher, accessed September 4, 2025, https://gregorykelleher.com/interview_practice_questions
8. C++ Atomic Tutorial - Complete Guide - GameDev Academy, accessed September 4, 2025, <https://gamedevacademy.org/c-atomic-tutorial-complete-guide/>
9. Atomics and Concurrency in C++ - freeCodeCamp, accessed September 4, 2025, <https://www.freecodecamp.org/news/atoms-and-concurrency-in-cpp/>
10. C++ 11 -
11. std::atomic - cppreference.com, accessed September 4, 2025, <https://en.cppreference.com/w/cpp/atomic/atomic.html>
12. Quickstart: Building with CMake | GoogleTest, accessed September 4, 2025, <http://google.github.io/googletest/quickstart-cmake.html>
13. Testing with Google Test - CMake-Episode 022 - YouTube, accessed September 4, 2025, <https://www.youtube.com/watch?v=W5jxjTsHXuA>